

January 16, 2024
 University of Waterloo, Ontario, Canada
 nilslukas.github.io / nlukas@uwaterloo.ca

Research Overview

I develop mechanisms to mitigate security and privacy threats that emerge when providing large machine learning systems with *untrustworthy* entities. These entities include (i) Untrustworthy Data that was manipulated to undermine the model's reliability, (ii) Untrustworthy Models that memorize training data and disclose private information to unauthorized users, and (iii) Untrustworthy Users who misuse the model's generated content to erode trust in digital media. To mitigate these threats, I study data poisoning and differential privacy and develop watermarking methods to detect generated content.

Education

Ph.D. in Computer Science , University of Waterloo, Canada Advisor: Prof. Florian Kerschbaum Thesis: "Analyzing Threats of Large-Scale Machine Learning Systems"	Jan 2019 - Feb 2024
M.Sc. in Computer Science , RWTH-Aachen, Germany Advisor: Prof. Stefan Decker Thesis: "Secure Inference of Deep Neural Networks"	May 2016 - Oct 2018
B.Sc. in Computer Science , RWTH-Aachen, Germany	Oct 2012 - Apr 2016

Experience

Research Intern , Microsoft Research Topic: Privacy for Large Language Models. See our IEEE S&P'23 publication.	May 2022 - Aug 2022
Instructional Advisor , University of Waterloo Course: Object-Oriented Software Development	Sep 2021 - Dec 2021
Co-Instructor , Group on Information Systems at the RWTH-Aachen Course: Data-Driven Medicine	Oct 2017 - Mar 2018

Conference Publications

[ICLR'24] AR: 31.0% (2250/7262)	Nils Lukas , Abdulrahman Diaa, Lucas Fenaux, Florian Kerschbaum. <i>Leveraging Optimization for Adaptive Attacks on Image Watermarks</i> . The Twelfth International Conference on Learning Representations, 2024 [pdf]
[ICLR'24] AR: 31.0% (2250/7262) 📰 Media Coverage	Benjamin Schneider, Nils Lukas , Florian Kerschbaum. <i>Universal Backdoor Attacks</i> . The Twelfth International Conference on Learning Representations, 2024, [pdf] [news].
[USENIX'23] AR: 29.2% (422/1444)	Nils Lukas and Florian Kerschbaum. <i>PTW: Pivotal Tuning Watermarking for Pre-Trained Image Generators</i> , 32nd USENIX Security Symposium, 2023, [pdf]
[IEEE S&P'23] AR: 17.0% (195/1147) 🏆 Distinguished Contribution Award at Microsoft MLADS	Nils Lukas , Ahmed Salem, Robert Sim, Shruti Tople, Lukas Wutschitz, Santiago Zanella-Béguelin. <i>Analyzing Leakage of Personally Identifiable Information in Language Models</i> , 44th IEEE Symposium on Security and Privacy, 2023, [pdf]
[IEEE S&P'22] AR: 14.5% (147/1012)	Nils Lukas , Edward Jiang, Xinda Li, Florian Kerschbaum. <i>SoK: How Robust is Image Classification Deep Neural Network Watermarking?</i> , 43rd IEEE Symposium on Security and Privacy, 2022, [pdf]
[ICLR'21] AR: 28.7% (860/2997) 📰 Spotlight (Notable 5%)	Nils Lukas , Yuxuan Zhang, Florian Kerschbaum. <i>Deep Neural Network Fingerprinting by Conferrable Adversarial Examples</i> , The Ninth International Conference on Learning Representations, 2021 [pdf]
[IH&MMSEC'21] AR: 40.0% (128/318)	Masoumeh Shafieinejad, Nils Lukas , Jiaqi Wang, Xinda Li, Florian Kerschbaum. <i>On the Robustness of Backdoor-based Watermarking in Deep Neural Networks</i> , Proceedings of the 2021 ACM Workshop on Information Hiding and Multimedia Security, 2021, [pdf]

[ACSAC'20]

AR: 20.9% (104/497)

Rasoul Mahdavi, Thomas Humphries, Bailey Kacsmar, Simeon Krastnikov, **Nils Lukas**, John Premkumar, Masoumeh Shafieinejad, Simon Oya, Florian Kerschbaum, Erik-Oliver Blass. *Practical Over-Threshold Multi-Party Private Set Intersection*, Annual Computer Security Applications Conference, 2020, [pdf]

[IEEE EuroS&P'20]

AR: 20.9% (39/187)

Bailey Kacsmar, Basit Khurram, **Nils Lukas**, Alexander Norton, Masoumeh Shafieinejad, Zhiwei Shang, Yaser Baseri, Maryam Sepehri, Simon Oya, Florian Kerschbaum. *Differentially private two-party set operations*, 2020 IEEE European Symposium on Security and Privacy (EuroS&P), 2020, [pdf]

Journal Publications**[AIP'18]**

Pascal Richter, Gregor Heimig, **Nils Lukas**, Martin Frank. *SunFlower: A new Solar Tower Simulation Method for use in Field Layout Optimization*, AIP Conference Proceedings, Volume 2033, Issue 1, 2018, [pdf]

Working Papers

2023

Nils Lukas, Florian Kerschbaum. *Pick your Poison: Undetectability versus Robustness in Data Poisoning Attacks against Deep Image Classification*, 2023, [pdf]

2023

Rasoul Mahdavi, **Nils Lukas**, Faezeh Ebrahimiaghazani, Thomas Humphries, Bailey Kacsmar, John Premkumar, Xinda Li, Simon Oya, Ehsan Amjadian, Florian Kerschbaum. *PEPSI: Practically Efficient Private Set Intersection in the Unbalanced Setting*, 2023, [pdf]

2023

Abdulrahman Diaa, Lucas Fenaux, Thomas Humphries, Marian Dietz, Faezeh Ebrahimiaghazani, Bailey Kacsmar, Xinda Li, **Nils Lukas**, Rasoul Akhavan Mahdavi, Simon Oya, Ehsan Amjadian, Florian Kerschbaum. *Fast and Private Inference of Deep Neural Networks by Co-designing Activation Functions*, 2023, [pdf]

Tutorials**[ESWC'18]**

Nils Lukas, Oya Beyan, Ali Hasnain. *Privacy-Preserving Information Extraction with Bloom Filters*, European Semantic Web Conference, Greece, 2018, [web]

Awards and Honors

2023

Best Poster Award - Presentation with David R. Cheriton [300 CAD]
Distinguished Contribution Award - Microsoft-internal 2023 MLADS conference

2022

David R. Cheriton Scholarship - 2-year award [20,000 CAD]
Outstanding Reviewer - International Conference on Machine Learning (ICML'22)

2019

Best Poster Award - Cybersecurity and Privacy Institute (CPI) [1,000 CAD]

2018

Excellence Graduation - Master's degree

2016

KU Global Scholarship - Semester abroad, Korea University [1.2 million KRW]

2014

MOGAM Scholarship - Semester abroad, Yonsei University, Korea [3,000 EUR]

University Service

2022

Student Board Member - Cybersecurity and Privacy Institute (CPI)
School Advisory Committee on Appointments (SACA) Liaison - CrySP lab

Academic Service**Program Chair**

RAID'24, CCS'23 (Artifact Evaluation)

Session Chair

IEEE S&P'23

External Reviewer

TheWebConf'24, ICLR'24, RAID'23, PETS'23, NeurIPS'23, PETS'22, NeurIPS'22, ICML'22 (*Top 10%* 🏆), PETS'21, CIKM'20

Teaching

2018 | *Data-Driven Medicine*, Co-instructor at the RWTH-Aachen in Germany

Invited Talks

Dec, 2023	<i>Analyzing Leakage of Personally Identifiable Information in Language Models</i> , Meta, [slides]
Oct, 2023	<i>How Reliable is Watermarking for Generative Machine Learning?</i> University of California, Berkeley, [slides]
June, 2023	<i>A Learnable Watermark for Deep Image Generators</i> Google, [slides]
June, 2023	<i>Analyzing Leakage of Personally Identifiable Information in Language Models</i> MongoDB, [slides]